

Aggregation and indexing are powerful features in MongoDB used to perform complex data processing and to optimize query performance.

Let's look at some examples using a sample `ecommerce` database with a `sales` collection that contains documents related to sales transactions.

## Step 1: Sample Data Setup

First, we'll create a sample dataset for aggregation and indexing examples.

javascript

Copy code

```
use ecommerce
```

```
db.sales.insertMany([
  { item: "laptop", quantity: 2, price: 800, customer: "Alice",
    region: "North", date: new Date("2024-10-01") },
  { item: "phone", quantity: 5, price: 300, customer: "Bob", region:
    "South", date: new Date("2024-10-02") },
  { item: "tablet", quantity: 3, price: 450, customer: "Alice",
    region: "North", date: new Date("2024-10-05") },
  { item: "laptop", quantity: 1, price: 800, customer: "Charlie",
    region: "East", date: new Date("2024-10-06") },
  { item: "phone", quantity: 2, price: 300, customer: "David",
    region: "West", date: new Date("2024-10-07") },
  { item: "laptop", quantity: 1, price: 800, customer: "Alice",
    region: "North", date: new Date("2024-10-10") }
])
```

## Step 2: Aggregation Queries

### 1. Total Sales Revenue by Item

This query calculates the total revenue for each item type.

javascript

Copy code

```
db.sales.aggregate([
  {
    $group: {
```

```

        _id: "$item",
        totalRevenue: { $sum: { $multiply: ["$quantity", "$price"]
    } }
    }
}
])

```

*Explanation:* This query groups documents by the `item` field and calculates `totalRevenue` by multiplying `quantity` and `price` for each document and summing up the result within each group.

## 2. Total Quantity Sold by Region

This query aggregates the total quantity of items sold per region.

```

javascript
Copy code
db.sales.aggregate([
  {
    $group: {
      _id: "$region",
      totalQuantity: { $sum: "$quantity" }
    }
  }
])

```

*Explanation:* This groups the documents by `region` and sums up the `quantity` field for each region.

## 3. Monthly Sales for Each Customer

Calculate the monthly sales for each customer to see their purchasing behavior.

```

javascript
Copy code
db.sales.aggregate([
  {
    $group: {
      _id: { customer: "$customer", month: { $month: "$date" }
    },

```

```

        totalSales: { $sum: { $multiply: ["$quantity", "$price"] }
      }
    }
  }
  })

```

*Explanation:* Here, the `$group` stage uses a compound key `_id` containing both `customer` and the month extracted from `date`. This gives the monthly sales for each customer.

#### 4. Average Purchase Amount by Item

This query calculates the average amount spent per item.

```

javascript
Copy code
db.sales.aggregate([
  {
    $group: {
      _id: "$item",
      averagePurchaseAmount: { $avg: { $multiply: ["$quantity",
"$price"] } }
    }
  }
])

```

### Step 3: Indexing

Indexes can significantly improve the performance of queries. Let's create indexes on fields used frequently in filtering, sorting, or grouping.

#### 1. Create an Index on the `item` Field

This index improves the performance of queries that filter or sort by the `item` field.

```

javascript
Copy code
db.sales.createIndex({ item: 1 })

```

#### 2. Create a Compound Index on `customer` and `date`

This compound index is useful for queries that involve both the `customer` and `date` fields.

javascript

Copy code

```
db.sales.createIndex({ customer: 1, date: -1 })
```

*Explanation:* The compound index supports queries that include both `customer` and `date`, such as looking up sales for a specific customer within a date range or sorting sales by date.

## Example Query Using the Index

After creating indexes, MongoDB will utilize them for queries involving these fields. Here's an example:

javascript

Copy code

```
db.sales.find({ customer: "Alice" }).sort({ date: -1 })
```

With the compound index on `{ customer: 1, date: -1 }`, this query will run faster than without the index, especially with a large dataset.

## Summary

- **Aggregation:** Useful for summarizing data. We used `$group` to calculate totals, averages, and monthly statistics.
- **Indexing:** Helps optimize query performance. We created single-field and compound indexes for frequently used fields in our queries.

Aggregation and indexing are essential for handling complex queries efficiently in MongoDB.